

## Introduction

Having multiple entities that all run simultaneously sharing data in some form or other, requires you to either design some scheme where protection is needless *OR* employ some form of *protection* and/or *signaling*. In this lecture we will look closely at the possible protection and signaling mechanisms available. This entails their concept and use.

## Contents

Theme(s)	Topic(s)	Reference
Thread Sync.	Protection and Synchronization	[3, chap. 30] [1]
	Posix semaphores (excluding named)	[3, chap. 53] [1]
Buffer case	Producer/Consumer impl. using semaphores	[2] [1]

## Material

### Slides

[1] S. Hansen, *Thread synchronization i*, Slides - see course repos.

### Online

[2] E. al. (). Producer-consumer problem. Wikipedia Article, [Online]. Available: [https://en.wikipedia.org/wiki/Producer%2dconsumer\\_problem](https://en.wikipedia.org/wiki/Producer%2dconsumer_problem).

### Hardback

[3] M. Kerrisk, *The Linux Programming Interface*. No Starch Press, Inc, 2010, ISBN: 978-1-59327-220-3.

## Fundamental questions to consider while reading

### Thread Synchronization I

- Shared data
  - What is shared data
  - What can and *will* go haywire with multiple threads - why
  - What to do about it
- Protection
  - Which different protection mechanisms are possible
  - What characterised each of these
  - How does the API look like and how do you use it

- Signalling
  - Which different signalling mechanisms are possible
  - What characterised each of these
  - How does the API look like and how do you use it
  - The notation *spurious wakeup* applies to *conditional variables* - how and what does this mean
- The Car park
  - How does it work conceptually
  - Determine who waits on who and when
- Producer/Consumer
  - Concept
  - From an implementation point of view, how does it work