

# Embedded Software

---

## Thread synchronization

# Agenda

---

- Why synchronization - Shared data problem revisited
- Cases
  - ▶ Sharing data between threads
  - ▶ The Producer / Consumer problem
  - ▶ Park-A-Lot 2000
- Types of synchronization methods

# The shared data problem revisited

# The shared data problem revisited

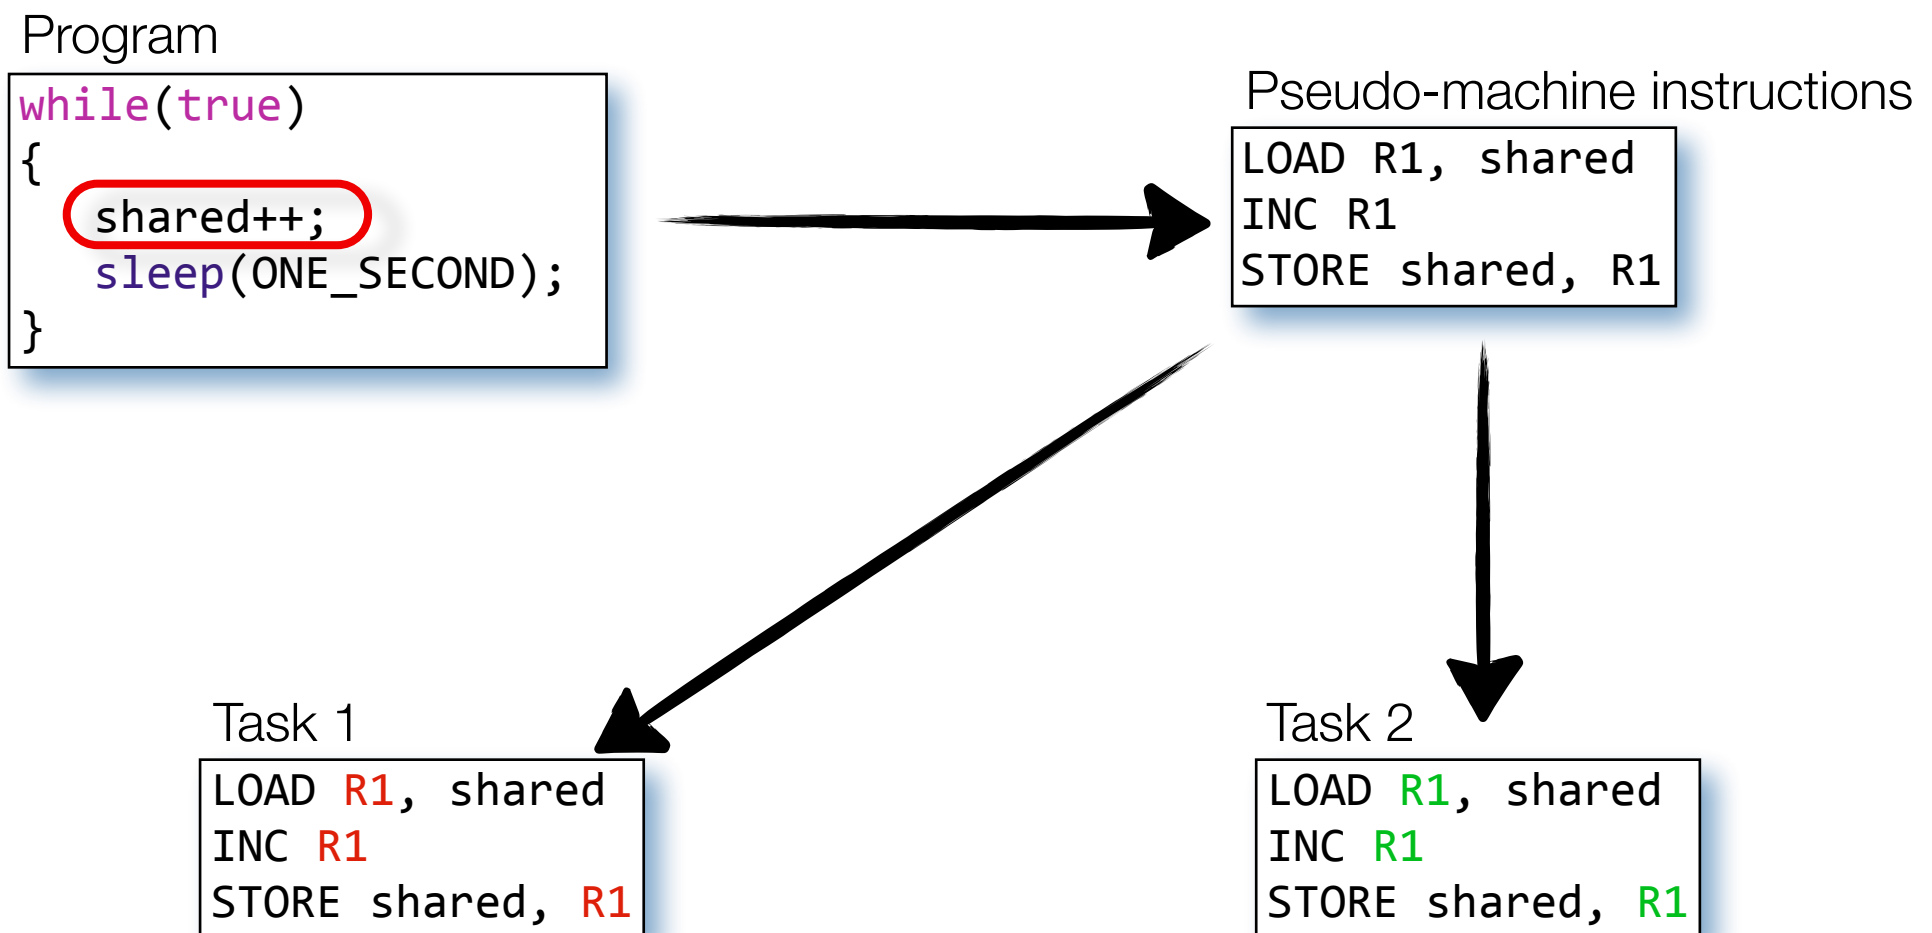
---

```
unsigned int shared;
void taskfunc()
{
    for(;;)
    {
        shared++;           // Increment i, then wait
        sleep(ONE_SECOND);  // 1 second
    }
}

int main()
{
    shared = 0;
    createThread(taskFunc); // Start two identical threads
    createThread(taskFunc); // that run the same function
    for(;;) sleep(ONE_SECOND); // 1 second
}
```

# The shared data problem revisited

- Let's zoom in:



# The Challenge

---

- We need a way to ensure that access to **shared** is mutually exclusive
  - ▶ When  $T_1$  is using **shared**,  $T_2$  must be denied access
  - ▶ When  $T_2$  is using **shared**,  $T_1$  must be denied access

# Cases

## Case - Ensuring shared data integrity




# Case - Ensuring shared data integrity

---

- Multiple thread entities read/write to a common data structure (maybe as simple as a single variable)

‣ ***Classic problem***



# Case - Ensuring shared data integrity

- Problem
  - ▶ Common *shared* variable
- Solution “a mutex”

```
unsigned int shared;
Mutex m = MUTEX_INITIALIZER;

void threadFunc()
{
    for(;;)
    {
        lock(m);
        shared++;
        unlock(m);
        sleep(ONE_SECOND);
    }
}

main()
{
    createThread(threadFunc); // Start two identical threads
    createThread(threadFunc); // that run the same function
    for(;;) sleep(100);
}
```

The mutex

Take the mutex  
(or block if needed)

// Increment i...

// then wait 1 second

Release the mutex

# Mutexes

---

- Mutexes are used to enforce MUTual EXclusion
- Mutexes are owned by one thread at a time - only the “**taker**” can release!
- Two operations on a mutex:
  - **lock(m)**
  - **unlock(m)**

```
lock(Mutex m)
{
    wait until m==1, then m=0;  // ATOMIC operation
}
```

If m=0, calling thread is **BLOCKED** until m==1  
If m==1, calling thread proceeds

```
unlock(Mutex m)
{
    m=1;  // ATOMIC operation
}
```

Now m==1 so a **BLOCKED** thread  
is made **READY**

# Case - Ensuring shared data integrity

- Problem
  - ▶ Common *shared* variable
- Solution “a semaphore”

```
unsigned int shared;
SEM_ID s;

void threadFunc()
{
    for(;;)
    {
        take(s);
        shared++;
        release(s);
        sleep(ONE_SECOND);
    }
}

main()
{
    s = createSem(1);
    createThread(threadFunc); // Start two identical threads
    createThread(threadFunc); // that run the same function
    for(;;) sleep(100);
}
```

Take the semaphore (or block if needed)

// Increment i...

// then wait 1 second

Release the semaphore

Initializing the sem to 1

# Semaphores

---

- Semaphores are used to enforce mutual exclusion or rather signaling
- Semaphores are NOT owned by one thread at a time - “**all**” can release!
- Two operations on a semaphore:
  - ▶ **take(s)** (A.K.A. get(s), pend(s), P(s), wait(s)...)
  - ▶ **release(s)** (A.K.A. give(s), post(s), V(s), signal(s)...)

```
take(Semaphore s)
{
    wait until s>0, then s=s-1;  // ATOMIC operation
}
```

If s=0, calling thread is **BLOCKED** until s>0  
If s>0, calling thread proceeds

```
release(Semaphore s)
{
    s=s+1;  // ATOMIC operation
}
```

Now s>0 so a **BLOCKED** thread  
is made **READY**

# Mutexes & Semaphores: FAQ

---


- “Can more than one thread wait for a mutex/semaphore at a time?”
  - ▶ Yes. The threads are queued – see next slide
- “Which of the blocked threads are made ready?”
  - ▶ FIFO: The thread that has waited the longest
  - ▶ Priority: The highest-priority thread

# Mutexes & Semaphores: The queue

- Each mutex/semaphore is associated with a waiting queue (FIFO/priority)
- When a thread takes a mutex:
  - ▶  $m=0$ : the next incoming thread is added to the mutex's queue
  - ▶  $m=1$ : running thread done, next thread activated

```
lock(Mutex m)
{
    wait until  $m==1$ , then  $m=0$ ;
}
```

```
unlock(Mutex m)
{
     $m=1$ ;
}
```




# The Producer-Consumer problem



# The Producer-Consumer problem

---


- A producer produces elements and *puts* them in a **buffer**, from which the consumer *retrieves* an element at a time



# The producer-consumer problem

---

- What happens if...
  - The producer **put()**'s into a full buffer?
  - The consumer **get()**'s from an empty buffer?
- How can this be handled?
  - Checking **insert** and **remove** before insertion?
    - ...and what if the buffer is full/empty? Sleep? How long?
  - Use 2 counting semaphores!



# The producer-consumer problem

- **Explain how to extend this implementation to use 2 counting semaphores to prevent buffer over/underrun**

```
template<typename T>
class Buffer
{
public:
    Buffer(size_t bufferSize) : buffer_(new T[bufferSize]),
        bufferSize_(bufferSize), insert_(0), remove_(0) { }

    void put(const T& x) {
        buffer_[insert_] = x;
        insert_ = (insert_+1)%bufferSize_;
    }

    T get() {
        T tmp = buffer_[remove_];
        remove_ = (remove_+1)%bufferSize_;
        return tmp;
    }

private:
    T* buffer_;
    size_t bufferSize_;
    CountingSemaphore emptySlotsLeft_;
    CountingSemaphore usedSlotsLeft_;
    size_t insert_;
    size_t remove_;
};
```

Implemented as **circular buffer**  
insert\_ : Insertion pointer  
remove\_ : Remove pointer

Only **one** producer & **one** consumer

Simplified construction  
Possible exceptions are NOT handled

# The producer-consumer problem

```
template<typename T>
class Buffer
{
public:
    Buffer(size_t bufferSize) : buffer_(new T[bufferSize]), bufferSize_(bufferSize),
                              insert_(0), remove_(0)
    {
        emptySlotsLeftSem_ = createCountingSem(bufferSize_);
        usedSlotsLeftSem_ = createCountingSem(0);
    }

    void put(const T& x) {
        take(emptySlotsLeftSem_);
        buffer_[insert_] = x;
        insert_ = (insert_ + 1) % bufferSize_;
        release(usedSlotsLeftSem_);
    }

    T get() {
        take(usedSlotsLeftSem_);
        T tmp = buffer_[remove_];
        remove_ = (remove_ + 1) % bufferSize_;
        release(emptySlotsLeftSem_);
    }

private:
    T* buffer_;
    size_t bufferSize_;
    SEM_ID emptySlotsLeftSem_;
    SEM_ID usedSlotsLeftSem_;
    size_t insert_;
    size_t remove_;
};
```

Semaphores are init with size  
bufferSize\_

The producer thread will automatically block  
if buffer is full on **put()**

The consumer thread will be auto-  
matically blocked if buffer is empty  
on **get()**

Both consumer and producer threads  
will be **automatically unblocked** if data is ready  
or buffer not full any longer


*Simplified construction  
Possible exceptions are NOT handled*

## Case - Park-a-lot 2000

# Case - Park-a-lot 2000

---

- Example: Park-a-lot 2000: An automated car parking system
  - ▶ One thread steers the car
  - ▶ Another thread steers the garage door opener
- ***Coordination how?***

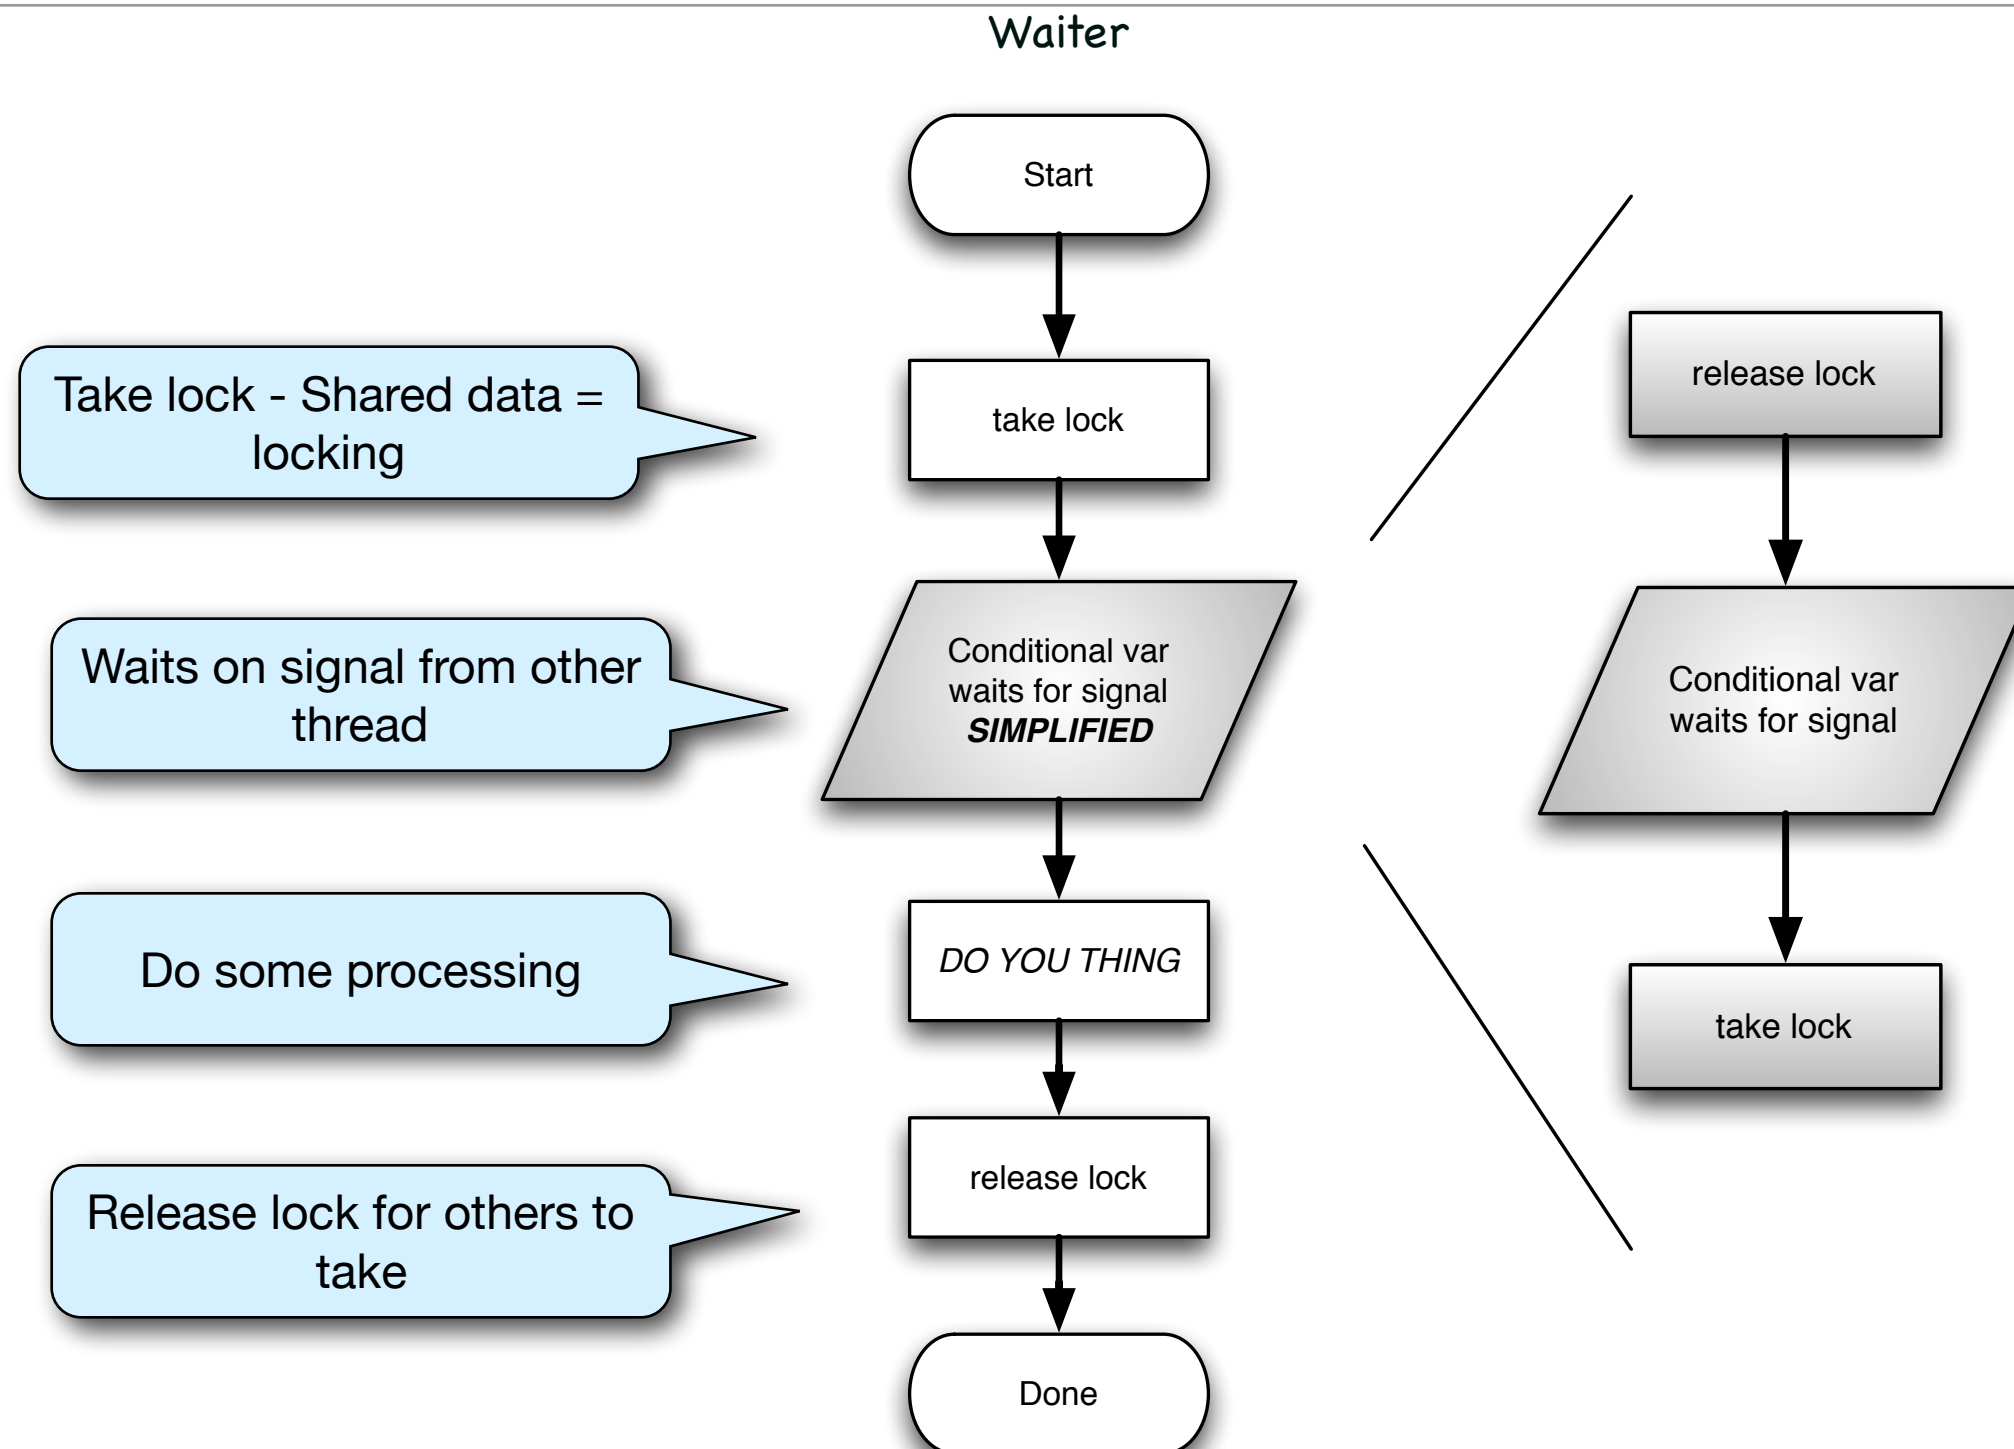


# Signaling mechanism - Which?

---

- We need Conditionals... But How?
  - ▶ Fundamental point is that we have a
    - ▶ Receiver/Waiter who waits on a conditional variable
    - ▶ Sender/Indicator who signals this particular conditional variable at some point

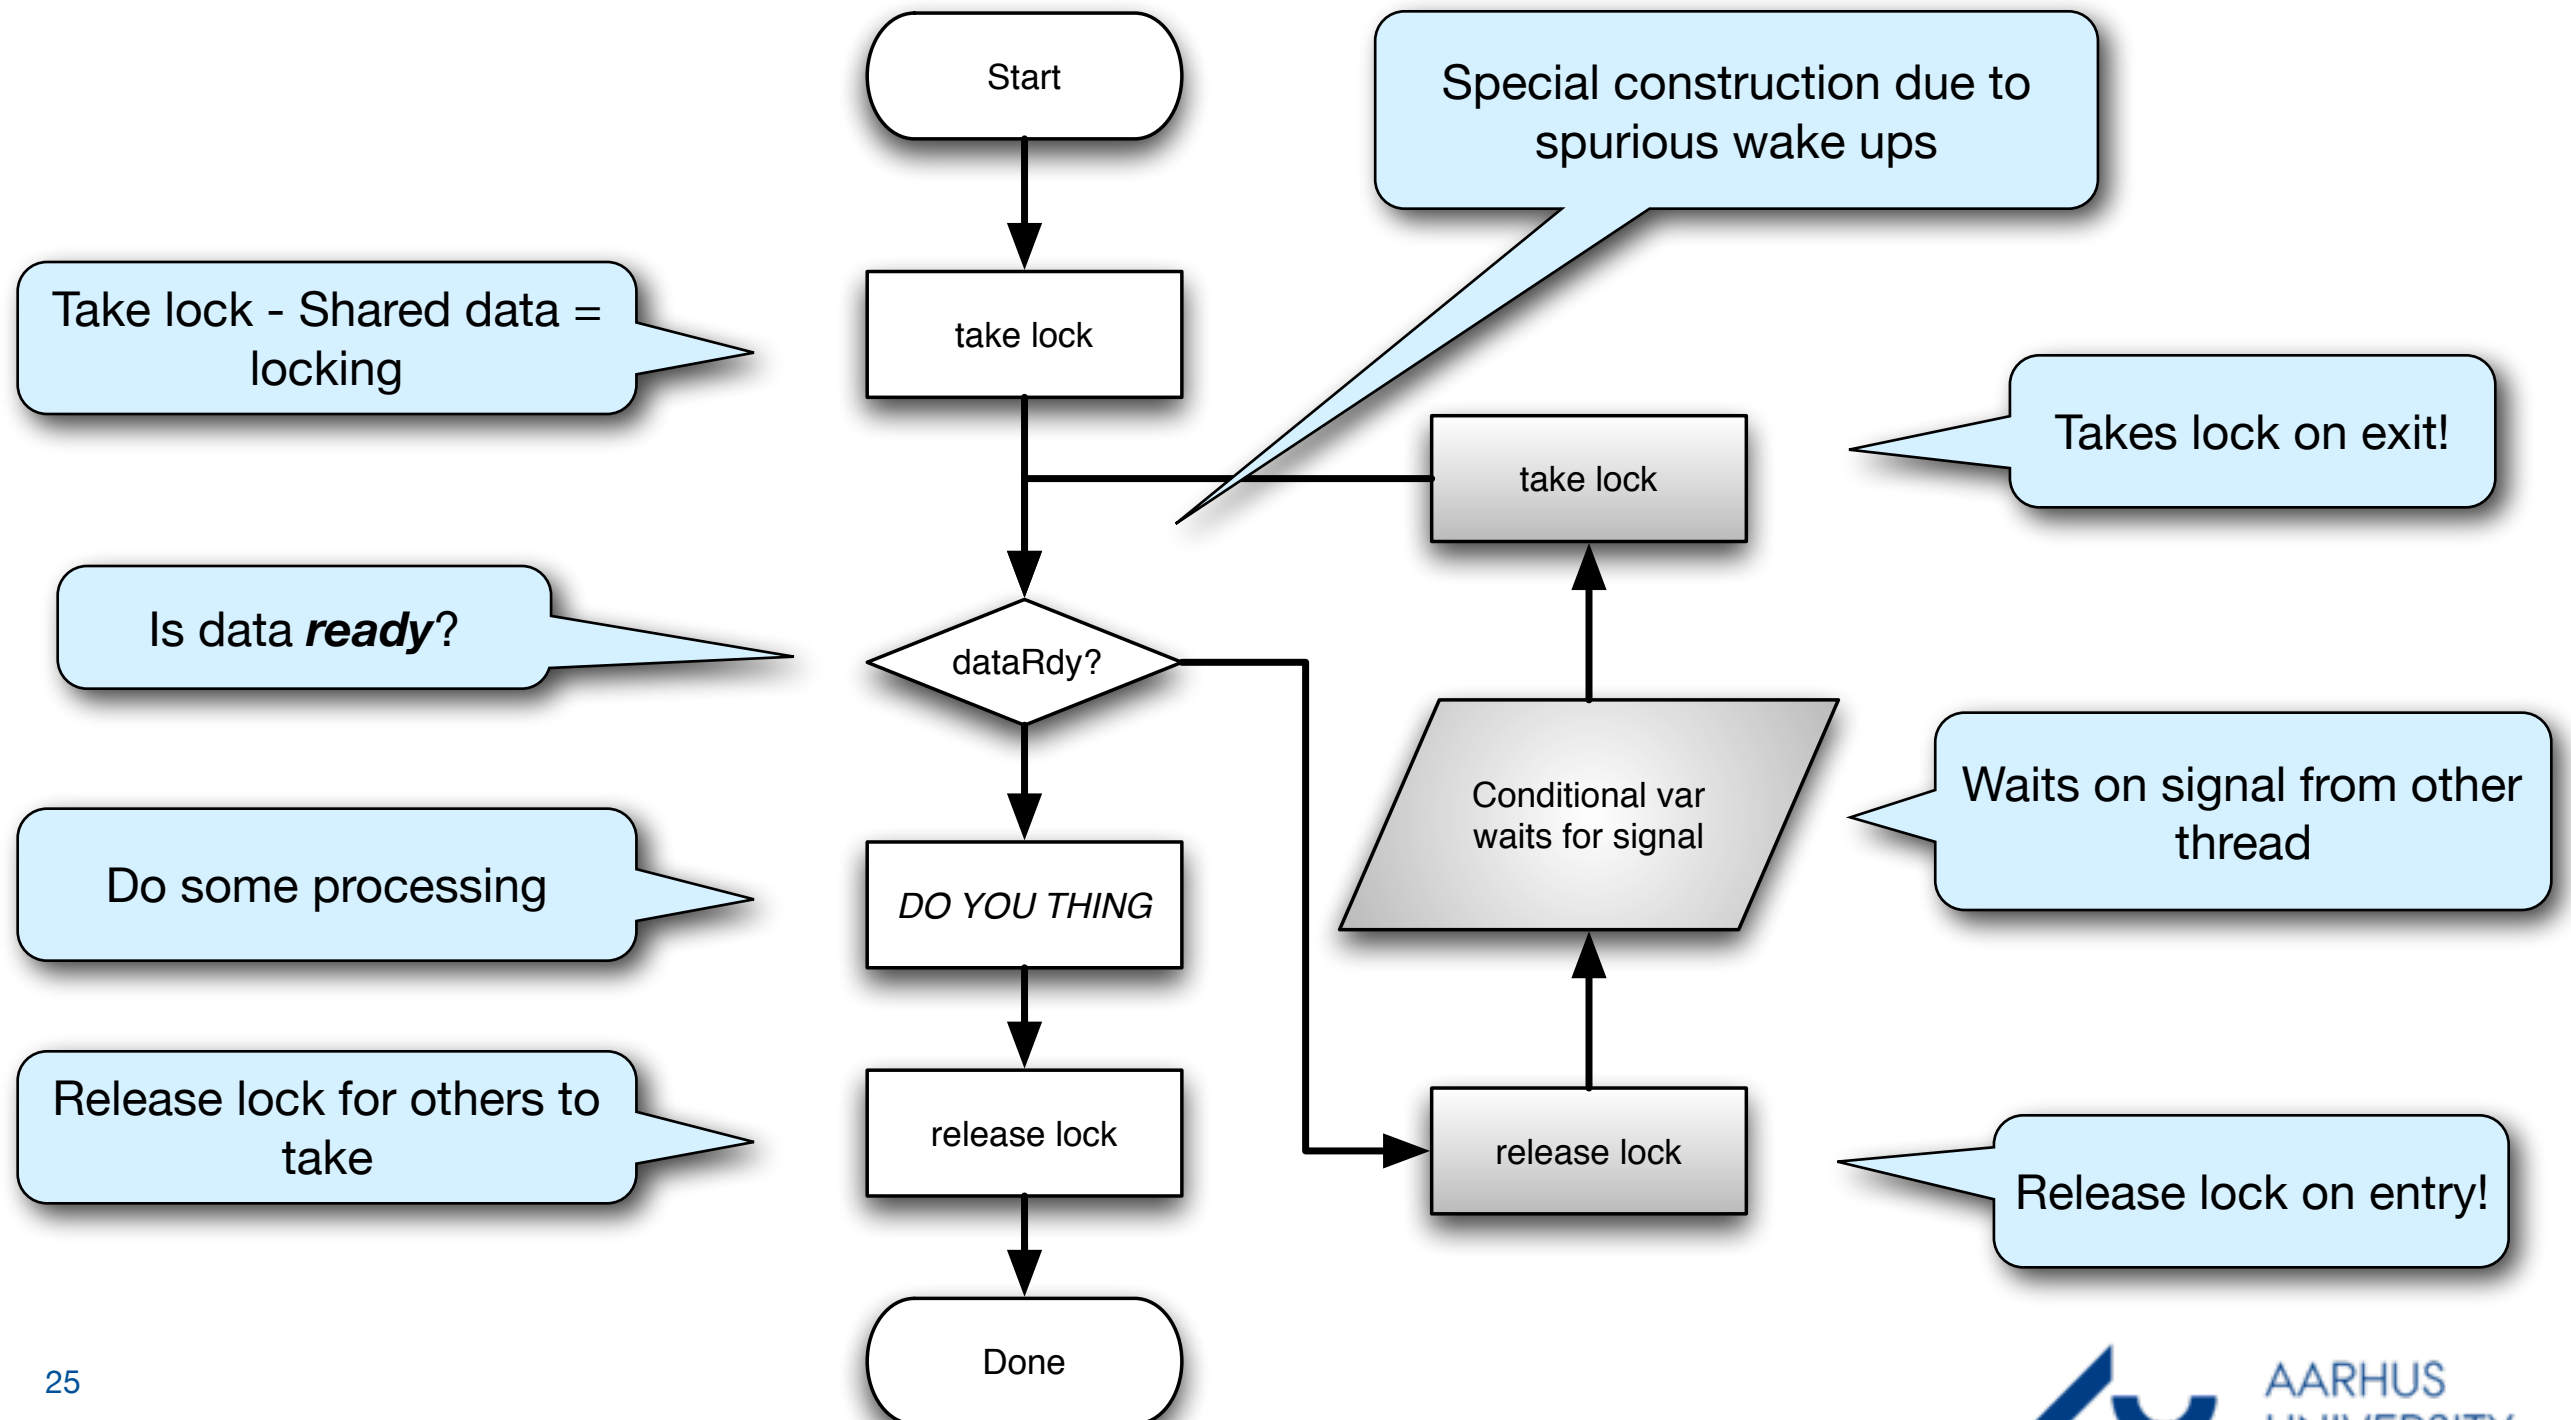
# We need Conditionals... - Receiver/Waiter



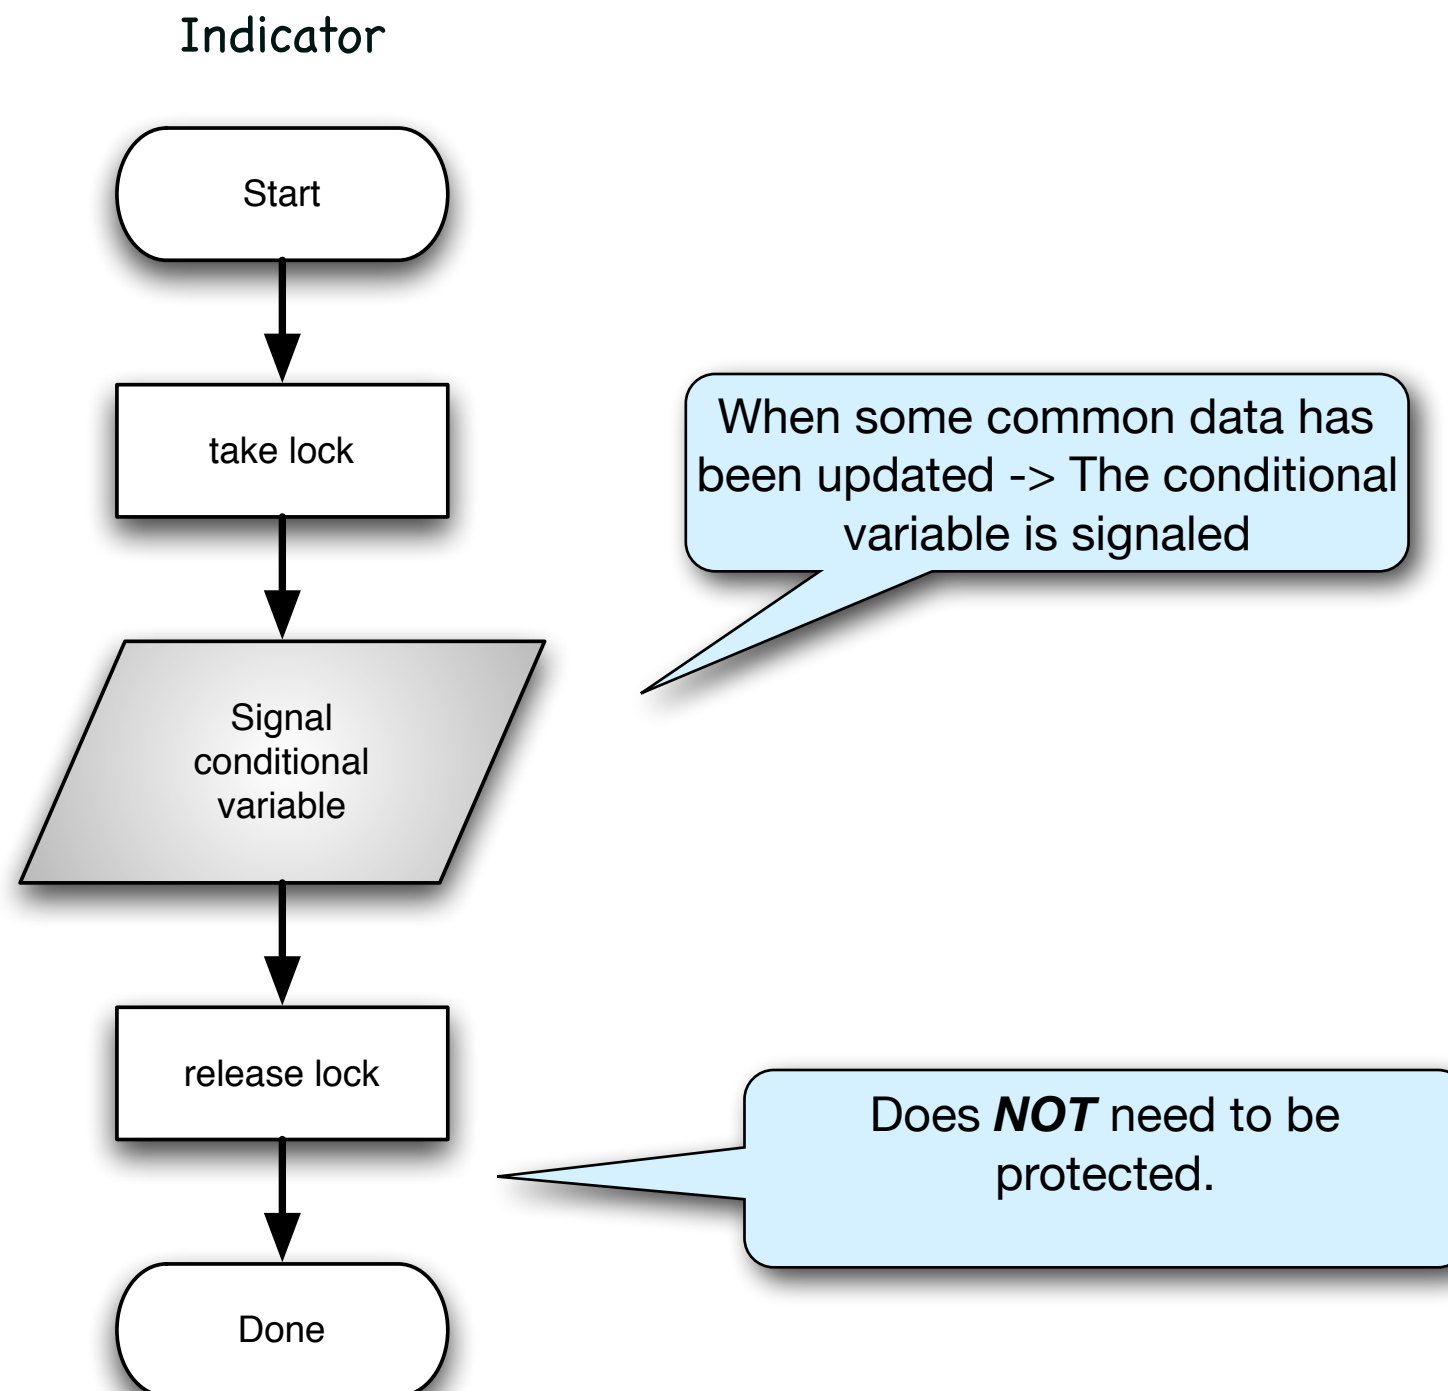


# We need Conditionals... - Receiver/Waiter

Waiter



# We need Conditionals... - Sender/Indicator



# Conditionals - How do you code it?

---

```
Mutex m;  
Conditional c;
```

```
void theWaiter()  
{  
    lock(m);  
  
    while (!what_we_are_waiting_for)  
    {  
        condWait(c, m);  
    }  
  
    unlock(m);  
}
```

```
void theIndicator()  
{  
    lock(m);  
    // Do something...  
    // unlock(m) - is okay  
    what_we_are_waiting_for = true;  
    condSignal(c);  
    unlock(m);  
}
```

# Park-a-lot 2000 - Feeble attempt

---

- Our first attempt:
- “hope”...another word system engineers don't like!

```
carDriverThread()  
{  
    driveUpToGarageDoor()  
    sleep(GARAGE_DOOR_OPEN_TIME);  
    // Let's hope the door is open!  
    driveIntoGarage();  
}
```

```
garageDoorControllerThread()  
{  
    openGarageDoor()  
    sleep(CAR_ENTER_GARAGE_TIME);  
    // Let's hope the car is in!  
    closeGarageDoor();  
}
```

- We need to be sure that...
  - ▶ The door is open before we move the car (car sync with garage door)
  - ▶ The car is in before we close the door (garage door sync with car)

# Park-a-lot 2000

- Our second attempt: Two-way synchronization

```
carDriverThread()  
{  
    driveUpToGarageDoor();  
    lock(mut);  
    carWaiting = true;  
    condSignal(entry);  
    while(!garageDoorOpen)  
        condWait(entry, mut);  
    driveIntoGarage();  
    carWaiting = false;  
    condSignal(entry);  
    unlock(mut);  
}
```

```
garageDoorControllerThread()  
{  
    lock(mut);  
    while(!carWaiting)  
        condWait(entry, mut);  
    openGarageDoor();  
    garageDoorOpen = true;  
    condSignal(entry);  
    while(carWaiting)  
        condWait(entry, mut);  
    closeGarageDoor();  
    garageDoorOpen = false;  
    unlock(mut);  
}
```

- This works!
  - ▶ 2-way synchronization
  - ▶ All waits are matched with signals

# Types of synchronization methods

# Types of synchronization methods

---

- Generally, there are three different types of semaphores for three different purposes:
  - ▶ Mutex:  $s=0$  or  $s=1$ , belongs to one thread at a time
  - ▶ Conditionals: Signaling facility used with a mutex
  - ▶ *Read/writable locks*: *Multiple readers - Exclusive writer*
  - ▶ Counting semaphore:  $s \geq 0$ , shared among threads
  - ▶ Binary semaphore:  $s=0$  or  $s=1$ , shared among threads

# POSIX Synchronization mechanisms (Not all included)

```
#include<pthread.h>

int pthread_mutex_init(pthread_mutex_t* mutex, pthread_mutex_attr_t *mutexattr);
int pthread_mutex_lock(pthread_mutex_t* mutex);
int pthread_mutex_unlock(pthread_mutex_t* mutex);
int pthread_mutex_destroy(pthread_mutex_t* mutex);

int pthread_rwlock_init(pthread_rwlock_t* mutex, pthread_rwlockattr_t *mutexattr);
int pthread_rwlock_rdlock(pthread_rwlock_t* mutex);
int pthread_rwlock_wrlock(pthread_rwlock_t* mutex);
int pthread_rwlock_unlock(pthread_rwlock_t* mutex);
int pthread_rwlock_destroy(pthread_rwlock_t* mutex);

int pthread_cond_init(pthread_cond_t *cond, const pthread_condattr_t *attr);
int pthread_cond_wait(pthread_cond_t *cond, pthread_mutex_t *mutex);
int pthread_cond_signal(pthread_cond_t *cond);
int pthread_cond_broadcast(pthread_cond_t *cond);
int pthread_cond_destroy(pthread_cond_t *cond)
```

```
#include<semaphore.h>


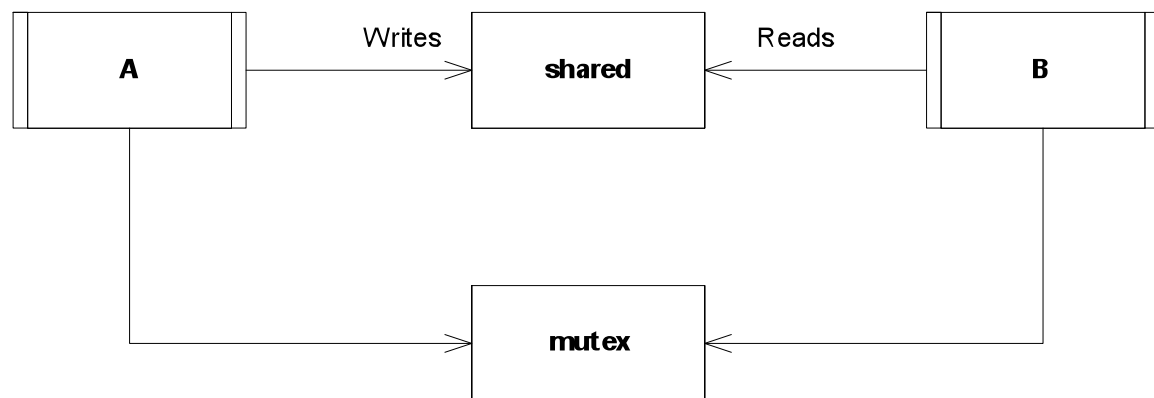
int sem_init(sem_t* sem,          // Semaphore ID
             int pshared,         // Unsupported. Set to 0
             unsigned int value  // Initial value , >=0
            );

int sem_destroy(sem_t* sem); // Semaphore ID to destroy
int sem_wait(sem_t* sem);   // Wait for sem
int sem_post(sem_t* sem);   // Post (signal) sem
```



# Aids / Tools

# The Monitor



- Monitor: A template class
  - ▶ When accessed, the Monitor
    1. takes mutex,
    2. accesses shared,
    3. releases mutex
  - ▶ Responsibility for mutual exclusion: Programmer → monitor
- Any drawbacks/consequences?
  - ▶ Complete copy of shared returned – takes time
  - ▶ Exception between lock() and unlock()?

# The Scoped Locking idiom

---

- A idiom pattern to ensure proper mutex clean-up, even on errors
- The idea: Create an object that automatically takes and releases a mutex at proper times – how?
  - ▶ lock() → constructor
  - ▶ unlock() → destructor
- ***How does this ensure clean-up?***
  - ▶ ***Generalized idiom called RAII - Learn IT!!!***